

Advanced NLP -

Session 4: Transformer Based Models

Prof. Dr. Richard Sieg
TH Köln IWS - WS 25/26



Agenda

01.

Tokenizer

Pre-Training

02.

Encoder Based & BERT

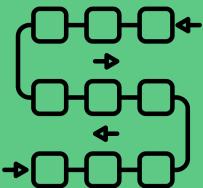
Tutorial

03.

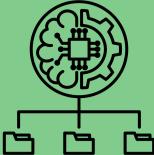
Decoder Based & GPT

Encoder-Decoder BART & T5

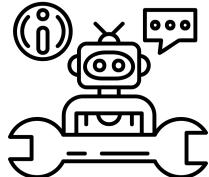
The 3 Ingredients of LLMs



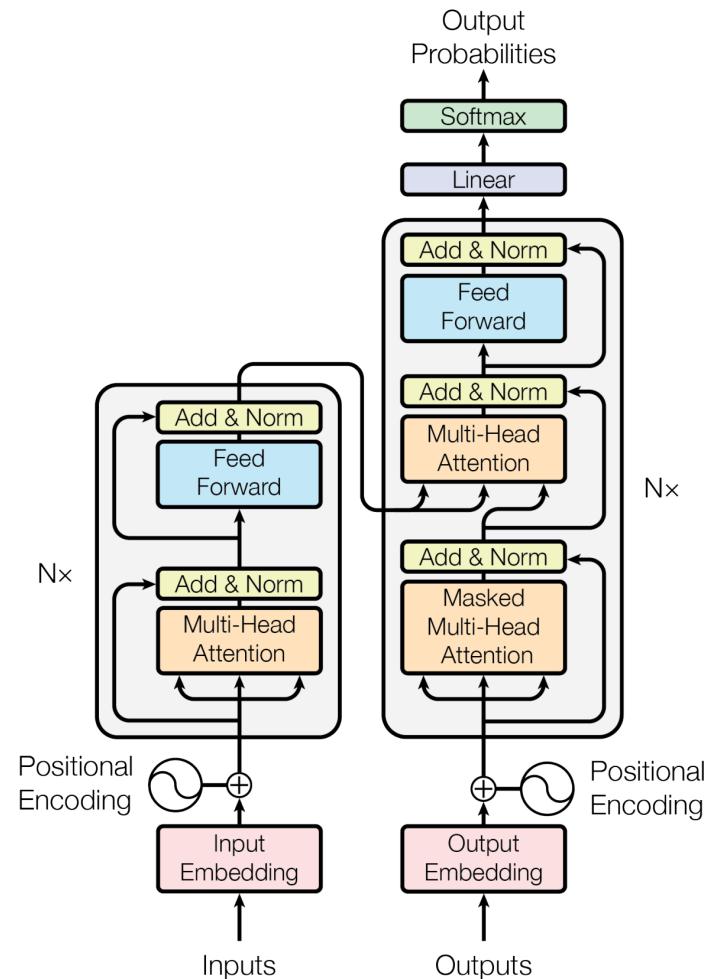
Process long sequences and context



Efficient training on huge datasets

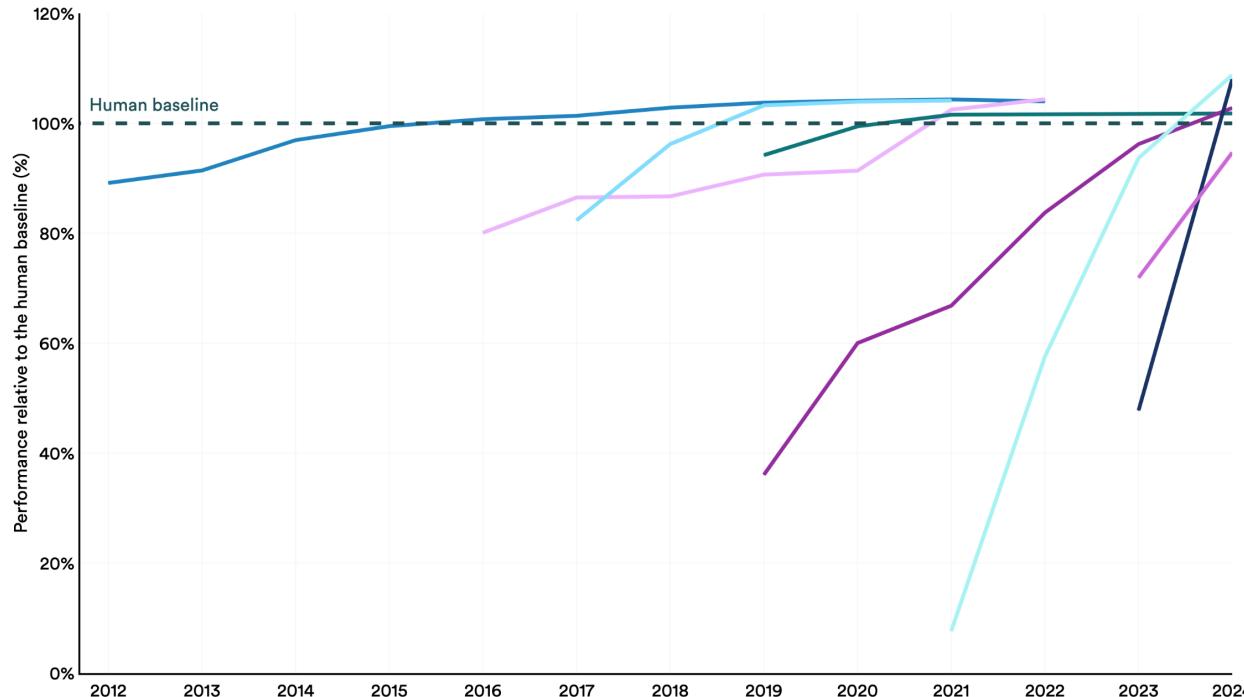


Follow (human) instructions



Select AI Index technical performance benchmarks vs. human performance

Source: AI Index, 2025 | Chart: 2025 AI Index report



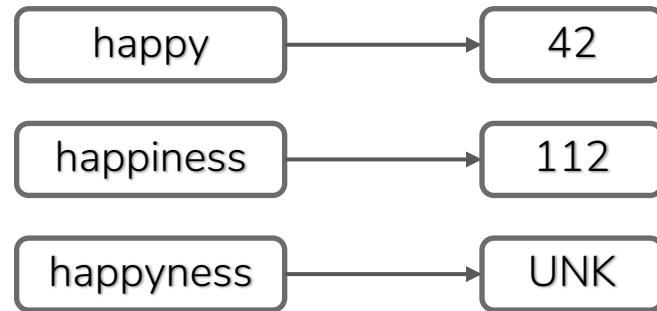
- Image classification (ImageNet Top-5)
- Medium-level reading comprehension (SQuAD 2.0)
- Multitask language understanding (MMLU)
- PhD-level science questions (GPQA Diamond)
- Visual reasoning (VQA)
- English language understanding (SuperGLUE)
- Competition-level mathematics (MATH)
- Multimodal understanding and reasoning (MMMU)

01.

Tokenizers

Fixed Vocabulary

- Up until now, we worked with a fixed vocabulary that we assign indices to
- Novel words are mapped to an <UNK> token
- **Issue 1:** We probably end up with a huge vocabulary
- **Issue 2:** We are unable to map novel words or leverage subword relations (e.g. “happiness” relates to “happy”)



Tokenizers

- **Token:** Any sequence of characters – mostly part of words (subwords)
- **Tokenization:** Splitting raw text into tokens

We <3 Python!

- Words: ["We", "<3", "Python", "!"]
- Substrings of length 3: ["We#", "<3#", "Pyt", "hon", "!"##"]
- BPE Tokenizer (GPT): ["We", "#<","3","#Python", "!"]

Byte Pair Encoding (BPE)

- **Goal:** Build a subword vocabulary up until a given size (e.g. 40k tokens)
- 1. Start with all ASCII characters and a symbol for “end of word”
- 2. We then check all pairs of characters and calculate their frequency in a corpus – the most common pairs are merged into new subwords
- 3. Replace the instances of these pairs with the subwords and continue with the algorithm until the subword vocabulary has the desired size
- This tokenizer algorithm is used for most LLMs as of today

WordPiece Tokenizer

- Tokenizer developed by Google for their BERT model
- Similar to BPE but merges not by frequency but a score that takes the frequency of the two parts in a merge into account

$$score(part_1, part_2) = \frac{freq(pair\ part_1 \& part_2)}{freq(part_1) \cdot freq(part_2)}$$

- This also takes into account what we loose by merging two subwords
- More difficult to implement and we might get out of vocabulary characters since we worked with predefined character set

02.

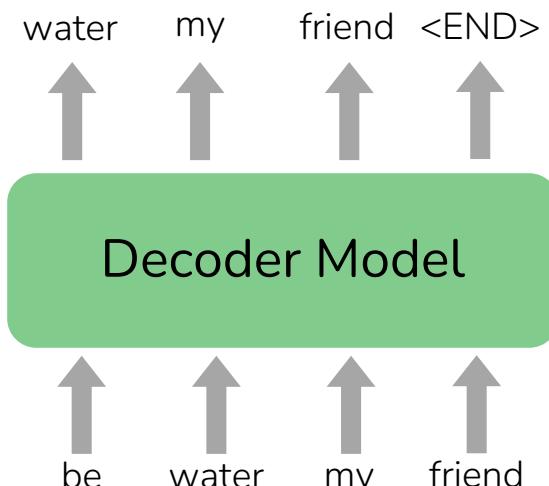
Pre-Training

From Pre-trained Word Embeddings to Pre-trained Models

- Up until 2017 we used pre-trained word embeddings as initial inputs for training an NLP task in an LSTM or Transformer
- These models are initialized with random parameters
- The models themselves have no initial idea of how language works and all contextual aspects of language
- Nowadays: The neural networks we use for a specific task are pre-trained on huge corpora
- Pre-trained models already have a strong representation of language and are a powerful initialization for neural networks in NLP
- The ImageNet moment of NLP

Pre-Training with Unsupervised Learning

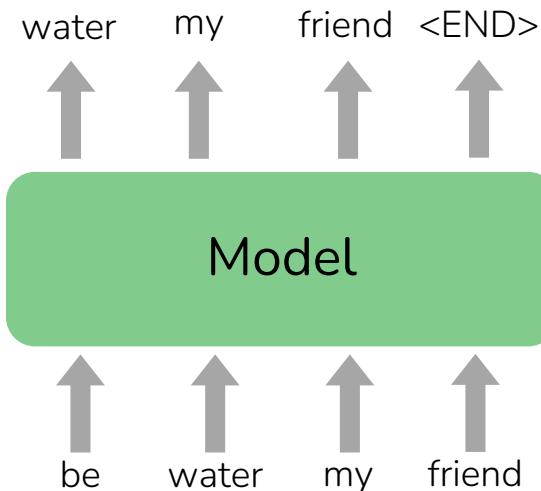
- Pre-Training a model works by masking some parts of the input and train the model to reconstruct these parts
- This allows us to process massive amounts of data without any labels (unsupervised)
- One option: Pretraining through language modeling, i.e. masking the next word and let the model predict the next word
→ Decoder Models (GPT etc)
- Or mask some random parts of the input like a cloze
→ Encoder Models (BERT etc)



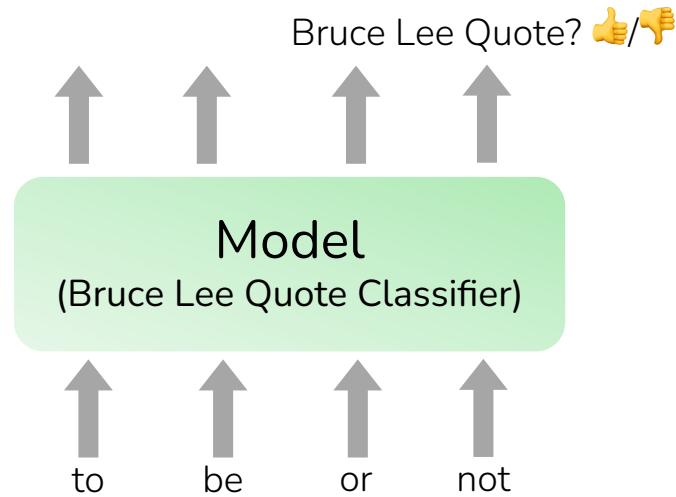
Pre-Training Paradigm

- Nowadays, we take a pre-trained model for our language (or a multilingual model) and use this as initial parameters to train our downstream task

Step1: Pre-train model on lots of text

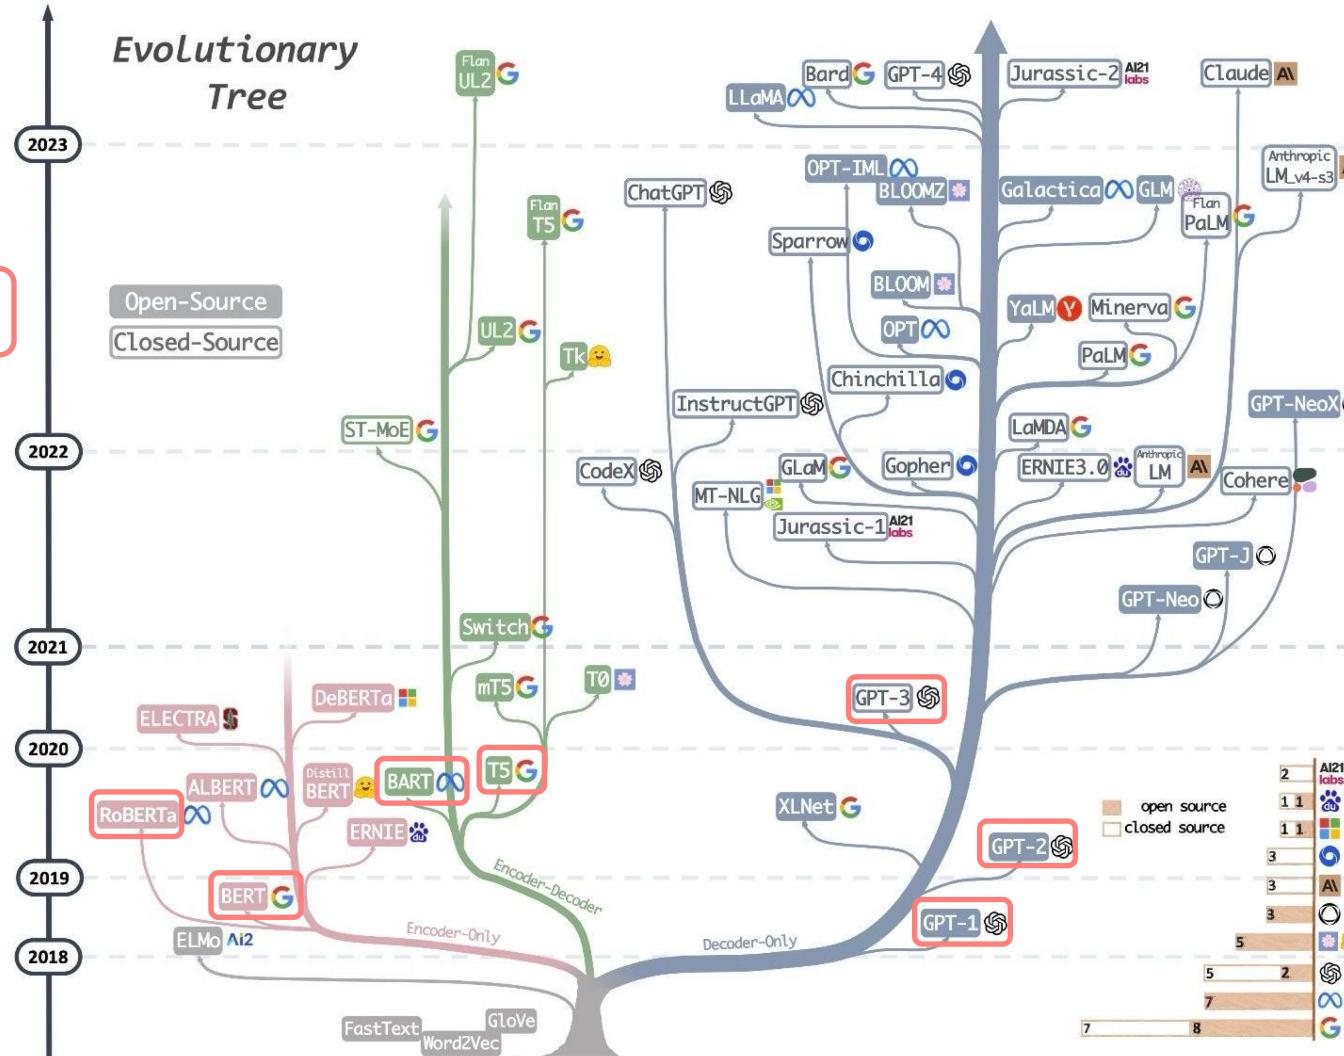


Step2: Fine-tune for your task

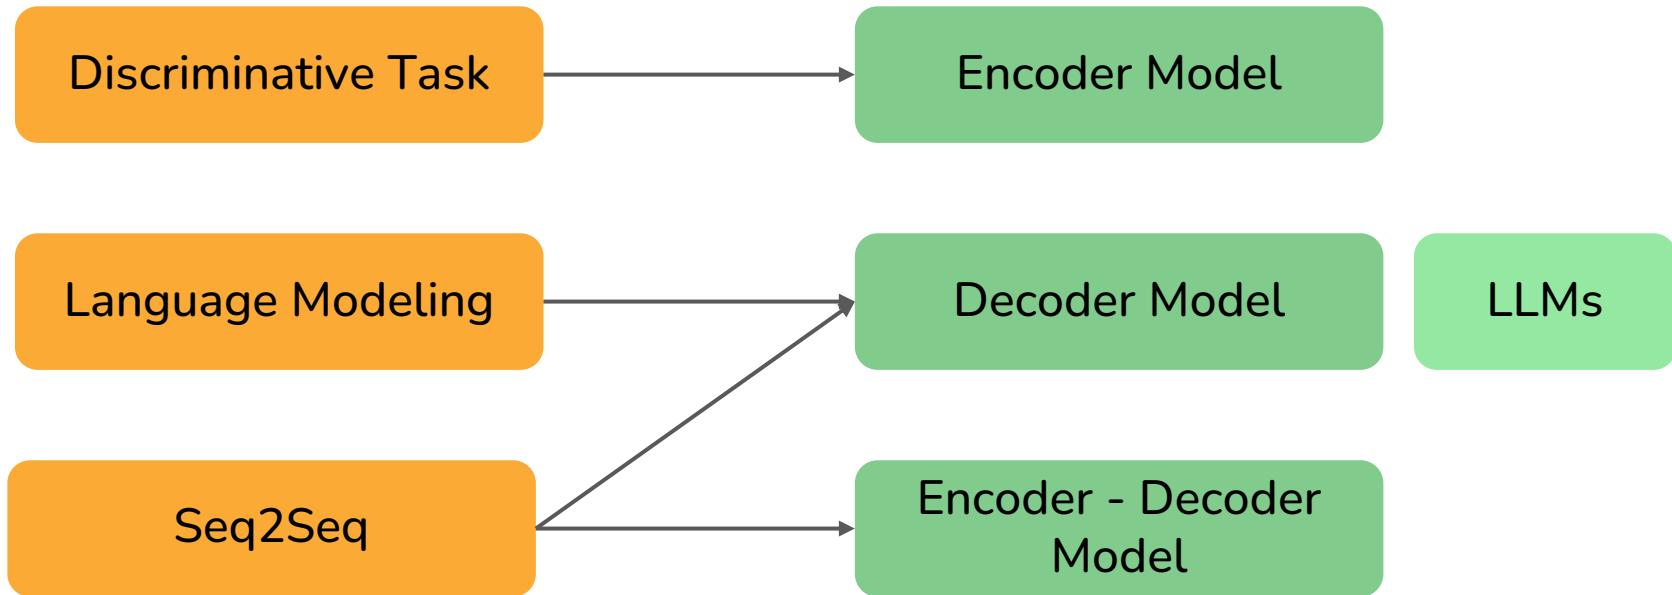


Evolutionary Tree

This lecture



Which Model for Which Task?

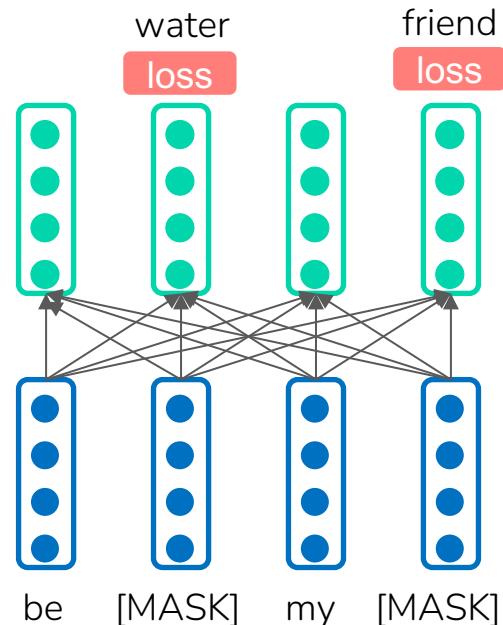


03.

Encoder Models &
BERT

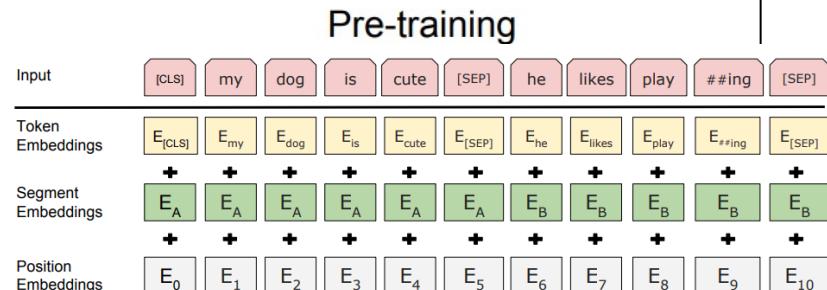
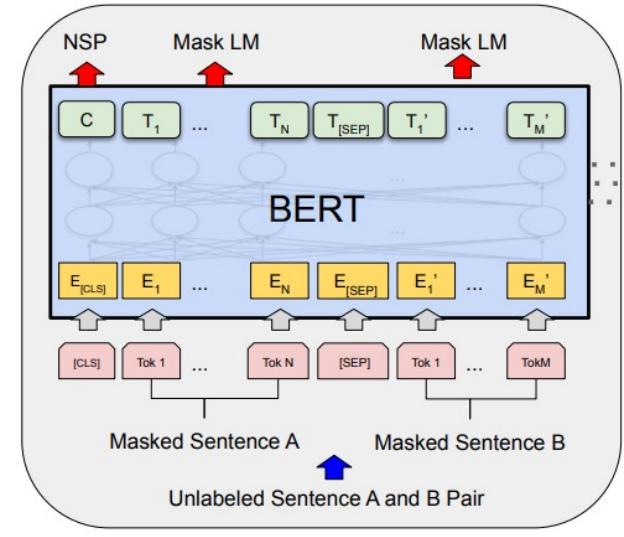
Pre-Training Encoders

- Encoders get bidirectional context and cannot be trained for language modeling
- **Masked Language Modeling:** Replace random words with [MASK] token and train the model to predict these words (like a cloze)
- We only compute the loss for the words that are masked out



BERT: Bidirectional Encoder Representations from Transformers

- Encoder based model developed by Google in 2018
- Trained on Masked LM and **next sentence prediction**
- Masked LM details:
 - 15% of tokens are masked and out of these
 - 80% are replaced with [MASK]
 - 10% replaced with random other token
 - 10% are unchanged but still predicted
- Next sentence prediction: predict whether one sentence follows the other or randomly sampled
 - Other BERT variants neglect this

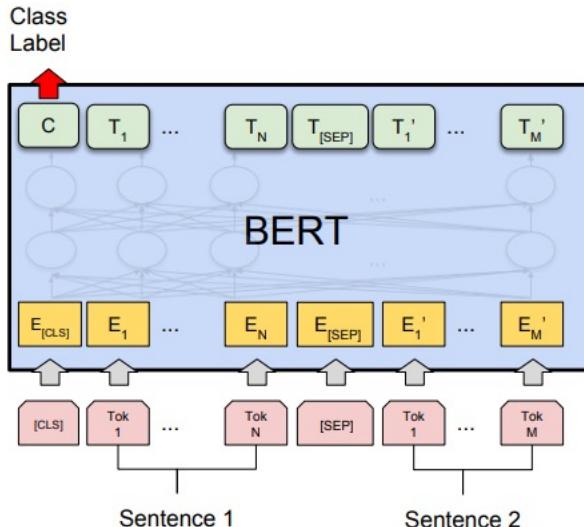


BERT Training Details

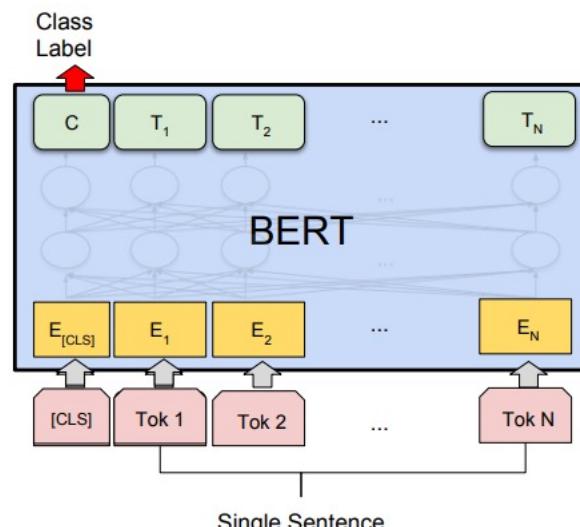
- BERT-base: 12 layers, 768-dim hidden states, 12 attention heads, 110 million params.
- BERT-large: 24 layers, 1024-dim hidden states, 16 attention heads, 340 million params.
- Trained on:
 - BooksCorpus (800 million words)
 - English Wikipedia (2,500 million words)
- BERT was pretrained with 64 TPU chips for a total of 4 days.

BERT Fine-Tuning Task – Text Classification

- For Text Classification, we start every sentence with a special [CLS] token
- The context vector of this token can then be used in a linear layer and softmaxed

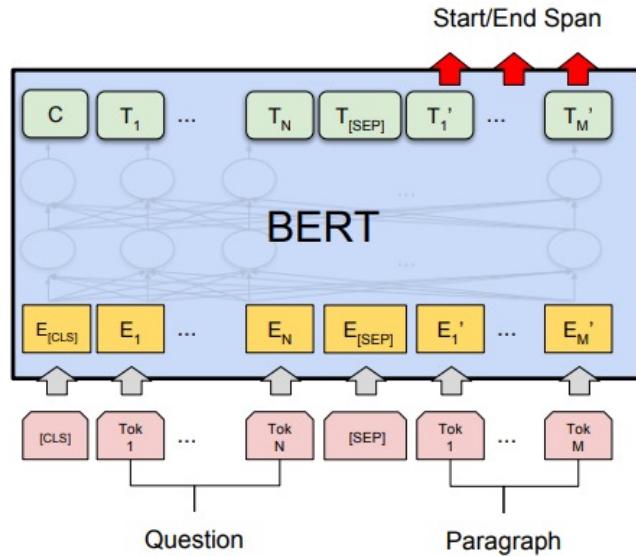


(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG

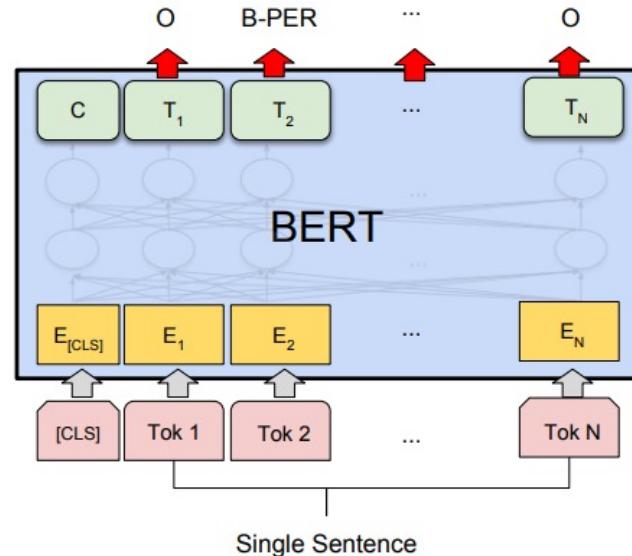


(b) Single Sentence Classification Tasks:
SST-2, CoLA

BERT Fine-Tuning Task – QA & Tagging



(c) Question Answering Tasks:
SQuAD v1.1



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

BERT Success

- By fine-tuning, BERT yielded state-of-the-art results on many tasks
- Is considered the powerhouse of modern NLP - in particular its variants
- Use it for any discriminative task in particular text classification

System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average -
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

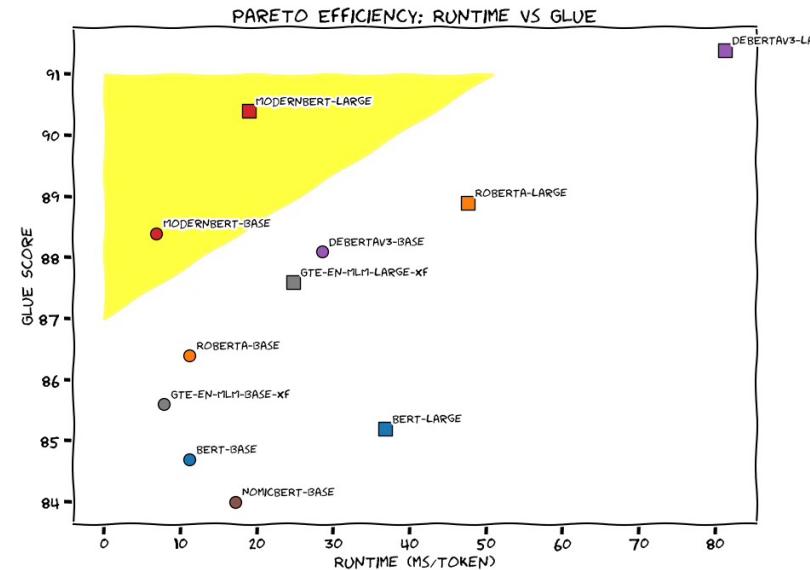
Table 1: GLUE Test results, scored by the evaluation server (<https://gluebenchmark.com/leaderboard>).

BERT Variants

- RoBERTa [2019]
 - No next sentence prediction
 - Dynamic masking (other masks every epoch)
 - BPE tokenizer instead of WordPiece
 - Small training adjustments
- DistilBERT
 - Much smaller version (40% less params) trained via student-teacher
- Lots of variants for different languages or use cases – just search HuggingFace 😊

modernBERT (12/2024)

- 8192 Tokens in the Context Window (vs 512 Vanilla BERT)
- Improved Transformer architecture using global and local Attention
- Lots of smaller tricks from years of LLM training experience
- EuroBERT (03/2025)
 - Basically modernBERT trained on European languages

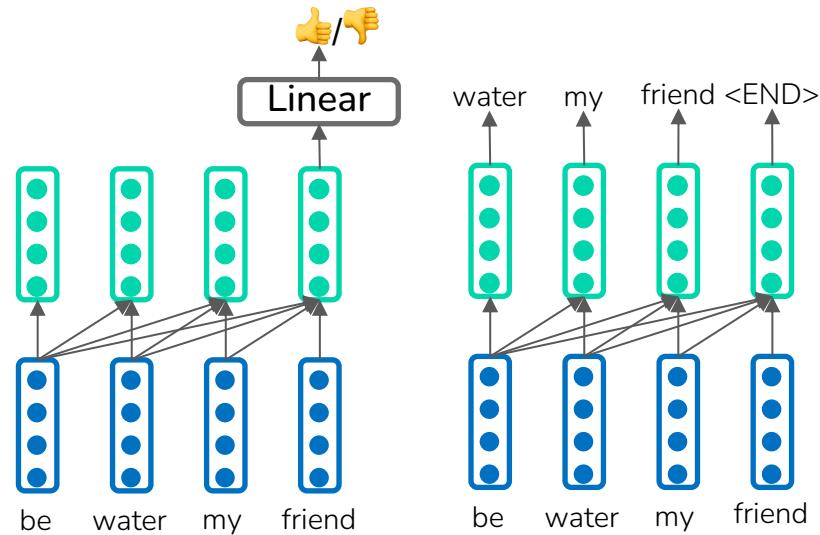


04.

Decoder Models & GPT

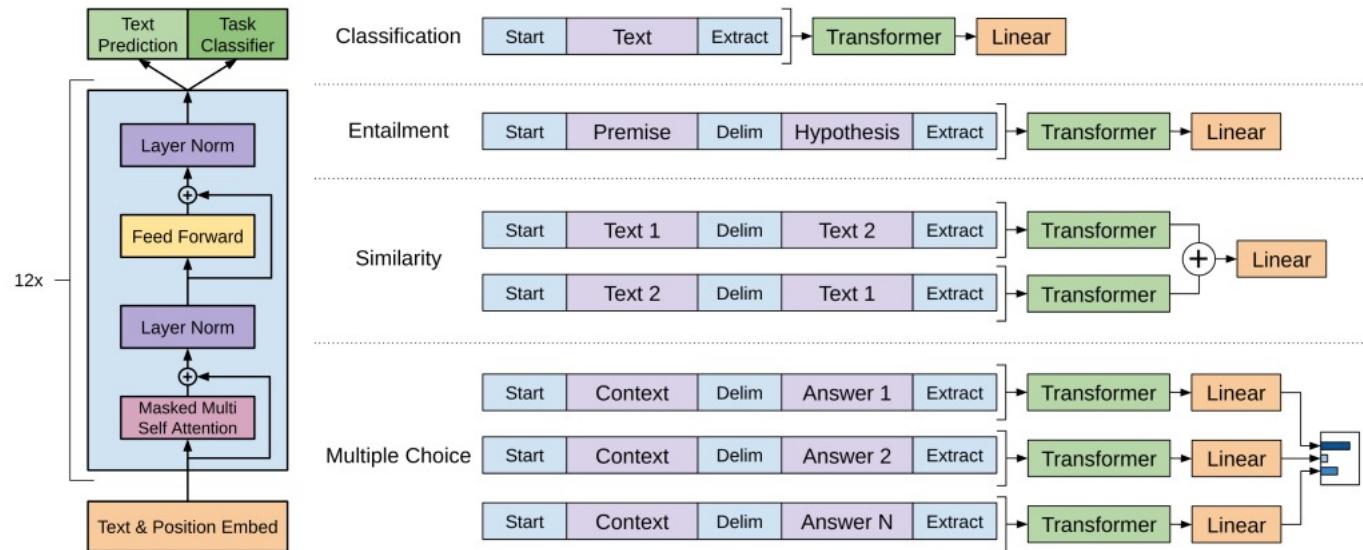
Pre-Training Decoders

- Decoders can also be fine-tuned for classification by adding a linear layer to the last hidden state (inferior to Encoder)
- More often: fine-tune them for generative seq2seq task, i.e. fine-tuning the language modeling task for a specific task



Generative Pretrained Transformer (GPT)

- Transformer Decoder with 12 layers
- Trained on BookCorpus
- BPE tokenizer of size 40k



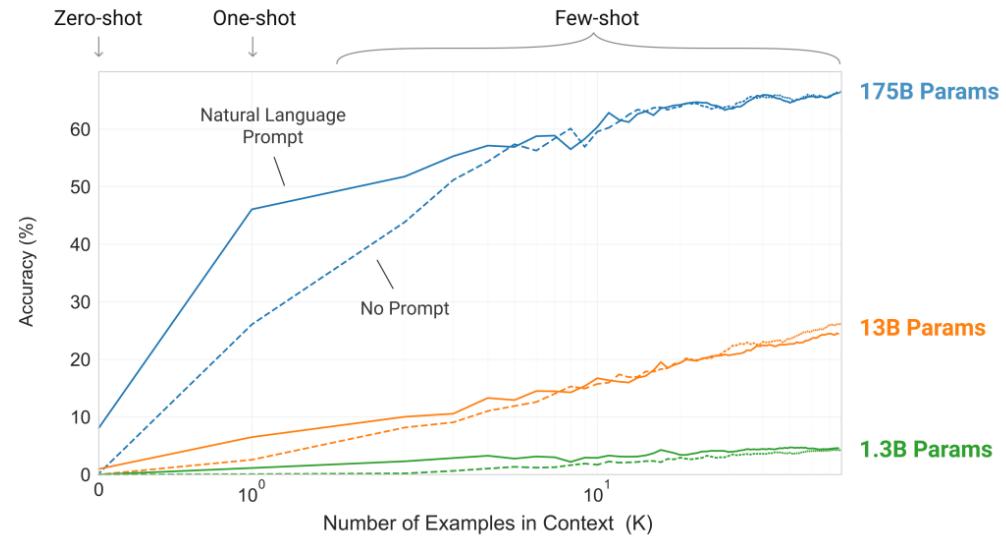
GPT2 [2019]

- Jump from 117m to 1.5b parameters and trained on much larger dataset (40Gb)
- First model that showed good results in ***zero-shot or few-shot prompting***
- Solve a task without fine-tuning, only prompt it with no or some examples
- And scaling seems increases the performance

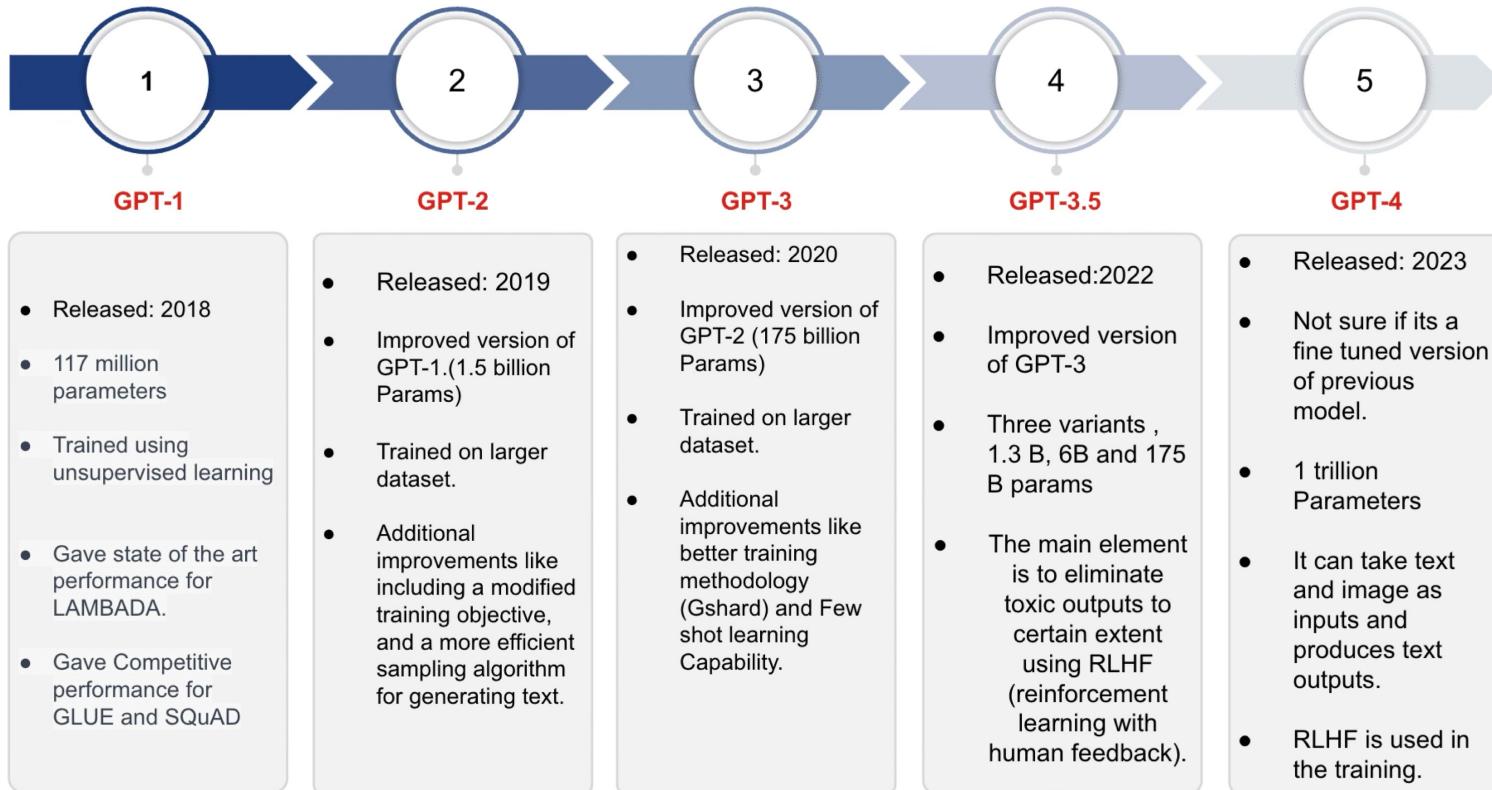
GPT3 [2020]

- Jump from 1.5b to 175b parameters and trained on 570Gb of data
- First model that performs exceptionally well with few-shot prompting
- ***The end of the fine-tuning paradigm!?***

Few-Shot Capabilities of GPT 3



GPT Development

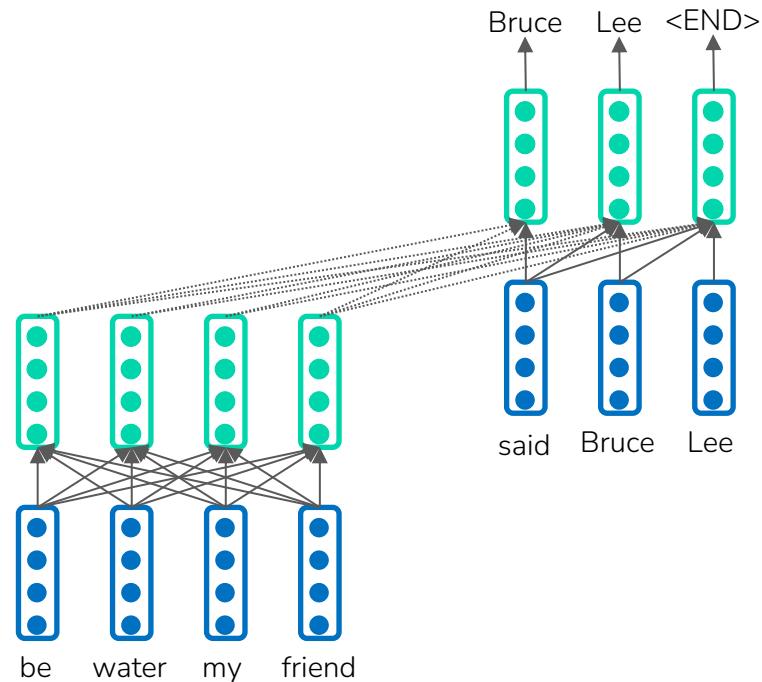


05.

Encoder Decoder Models

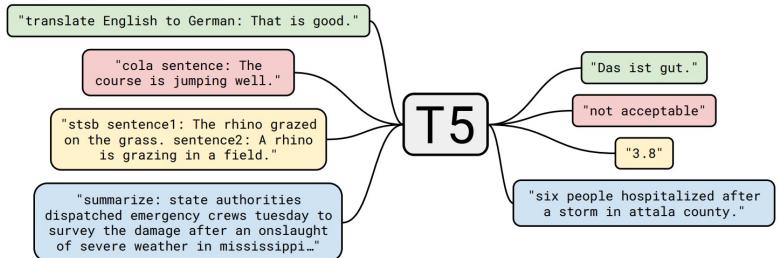
Pre-Training Encoder-Decoder

- Train like language modeling but give the first half of a sequence to the encoder and do not predict it and the decoder learns to predict the second half of the sequence

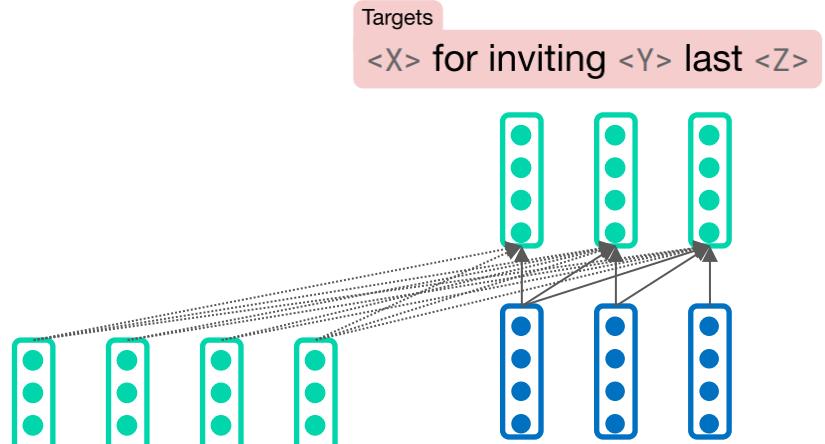


T5 [Google, 2019]

- Text-to-Text Transfer Transformer (T5)
- Trained on **span corruption** task
- Replace spans of different length from the input with placeholders and learn to fill the spans

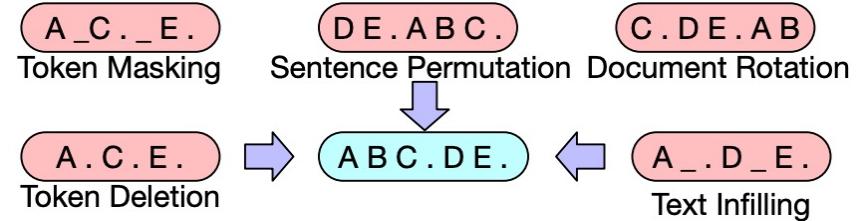
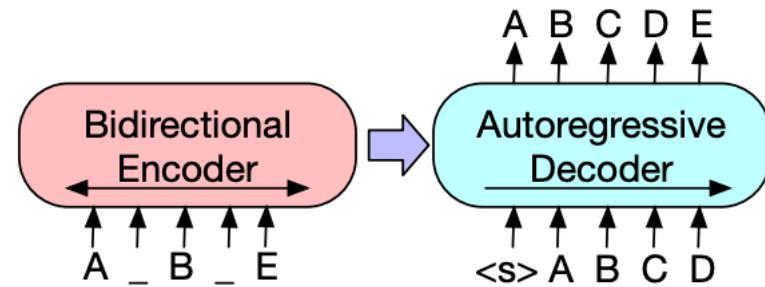


Original text
Thank you for inviting me to your party last week.
Inputs
Thank you <X> me to your party <Y> week.

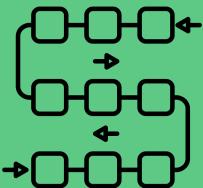


BART [Meta / FAI, 2019]

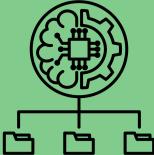
- Mixes BERT and GPT training
- Adds noise to documents in multiple ways and then learns to reconstruct the original document (**denoising**)



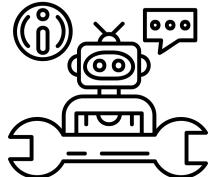
The 3 Ingredients of LLMs



Process long sequences and context



Efficient training on huge datasets



Follow (human) instructions

Tutorial

